

Module Five

CSG1105 Applied Communications

Functions of the Transport Layer

- Provides Application multiplexing (ports)
- Provides end-to-end delivery of data
- Segments messages (TCP)
- Ensures segments received in the correct sequence (TCP)

Application Multiplexing

- Any process requiring a TCP/IP connection identifies itself via a 16 bit number known as a **port address** or **port**
- Ports are one of two categories:
 - ▶ Well-known Ports: numbers 1-1023* allocated to servers e.g. 22 - SSH, 25 - SMTP, 80 - HTTP. These ports are allocated by the IANA (Internet Assigned number Authority)
 - ▶ Ephemeral Ports: numbers 1024 - 65535, generally used by client process and are assigned as needed
- There are many servers that use port numbers above 1023, some of which are mapped by the IANA, some of which are accepted by general usage
- The IP address Port number combination becomes a unique network address for an application
- The combination of each of the client and server IP Address/Port number pairs identifies a connection and is used in the most common TCP/IP API the **Socket**

- A **Socket** is a programming API for TCP/IP introduced in 4.2BSD Unix (Berkeley Software Distribution) first introduced in 1983
- The socket is the *de facto* industry standard for accessing TCP/IP from an application
- Sockets use a **file handle** to access network services as the connection is treated as a **byte stream**
- A socket address consists of **protocol, IP address, port** e.g.
`tcp, 192.168.0.50, 80`
- A connection is represented by an *association* which has the IP Address/Port of both client and server e.g..
`tcp, 192.168.0.50, 80, 172.16.100.22, 50111`
- Each connection on a host may be uniquely identified by this combination

- The Transport Layer provides two protocols
- ① UDP: User Datagram Protocol - single packet with Transport Layer header
- ② TCP: Transmission Control Protocol - reliable byte stream

- UDP, the User Datagram Protocol, is an application interface to IP that provides port addresses for a single packet
- UDP uses *ports* as previously described for connection multiplexing
- UDP is *connectionless*: it does not provide any reliability, flow control or error recovery which must be handled by the application if required.
- UDP has far less overhead for transmission than TCP and is used in some streaming applications for this reason
- UDP is also used for services like DNS,TFTP,RPC & LDAP, that only require small amounts of data
- The maximum payload size of a UDP Datagram is 516 bytes

- TCP is a *connection-oriented* protocol which provides a *full-duplex* byte stream between processes
 - ▶ TCP uses *ports* as previously described for connection multiplexing
 - ▶ The application writes and reads a *byte stream* to/from the TCP connection
 - ▶ TCP *Segments* the byte stream to be passed to the Internetwork layer (done with knowledge of the Internetwork layer MTU)
 - ▶ TCP provides *Reliability* using sequence numbers and acknowledgements (ACK)
 - ▶ TCP provides *Flow control* via the receiver ACK mechanism (*sliding window*)

- As previously mentioned, TCP is *connection oriented*. Think of it like a phone call. . .
 - 1 You enter a phone number you wish to call and the remote handset either rings or you get an engaged signal
 - 2 If the is engaged, the call attempt ends
 - 3 If the phone rings and someone answers, a brief dialogue occurs ensuring you're talking to the correct person
 - 4 Phone conversation continues until completion if no network errors
 - 5 Caller or called ends call
- Phone calls are connection oriented

- TCP uses a process called the *Three Way Handshake* to establish a connection
 - 1 The server must have a process listening for incoming connection requests it creates a Socket with a *Passive Open*
 - 2 A client attempting to connect to the server creates an *Active Open* call
 - a. The Client sends a SYN (synchronise) message with a random sequence number (SEQ), say 2044
 - b. If not accepting requests, the server can reply with a NAK (negative acknowledge)
 - c. If accepting requests, the server replies with a SYN ACK (acknowledge) with it's own random SEQ, say 42000, and acknowledges the request with the next expected SEQ, 2045
 - d. The client completes the connection with an ACK with SEQ 2045 and ACK 42001

- ③ The connection is now open and segments may be exchanged using the sequence numbers established
- ④ Data is exchanged and the sequence numbers are used for reliability and flow control
- ⑤ The connection is closed with one end sending a FIN segment , the other end will complete any operations in progress then sends a FIN segment.
- ⑥ When closed, all status fields etc are cleared and if appropriate an new passive open is performed on the server to again listen for incoming connections.

In practice, a busy server such as a web server, will have multiple listening processes that can be used to establish a connection.

- Each TCP connection has a *sliding window* that assists in reliability and flow control
- The “window” is a range of unacknowledged segments that may be transmitted
- If the window is exceeded, the sender ceases sending until an appropriate ACK of sent segments is received
- If a receiving end is congested, it stops sending ACKs which will shut down the sender
- if a segment is lost, the receiving end will repeat sending ACKs for the last correctly received segment until the sender retransmits the lost segment
- The receiver, if it has segments after the lost segment still in its buffer, can acknowledge multiple segments when the lost segment is received.

- Now that we have a basic understanding of TCP and IP addresses, we can explain a commonly used feature called **Network Address Translation**.
- There are two forms under this banner
 - 1 Network Address Translation (NAT)
 - 2 Port Address Translation (PAT - AKA Network Address Port Translation NAPT)
- PAT is now the most commonly used of the two, but understanding the basics of NAT will assist in the comprehension of PAT

Interlude: Private IP Addresses

- The issue of address exhaustion and the use of IP in non-Internet connected networks resulted in the publication of RFC1918 - Address Allocation for Private Internets
- This document outlined the available addresses and appropriate usage of three ranges of IP addresses
- It also mandated that these not be used on the wider Internet

Start	Finish	Mask	Net Size
10.0.0.0	10.255.255.255	/8	16777214
172.16.0.0	172.31.255.255	/12	1048574
192.168.0.0	192.168.255.255	/16	65534

- Internal use of a large space allows arbitrary codes to be embedded in addresses
- At ECU: **10.18.2.51/24** - Building 18, Level Two allowing 254 host per level

- Private IPs cannot be routed on the wider Internet
- NAT was originally meant to be a stopgap solution to IP address exhaustion while waiting for IPv6
- It also allowed organisations previously not connected to the Internet, but using IP internally to connect without having to renumber all networks
- Refer to RFC3022: Traditional IP Network Address Translator for details

- A NAT gateway (usually a router), maintains a table of available external IP addresses
- When an internal host wants to establish a connection to an external host, the NAT gateway assigns an external IP to the connection and **rewrites** the IP header of the outgoing packet with the external IP
- A table is maintained with a tuple of the internal/external IP addresses
- When a packet is received for the internal host from the external host, the table is referenced and the IP header is again re-written
- When the connection is closed, the entry is removed from the table

Host	Destination	Source
Internal Outgoing	134.1.1.99:80	10.18.2.51:446
Gateway Outgoing	134.1.1.99:80	228.50.21.1:446
External Received	134.1.1.99:80	228.50.21.1:446
External Sent	228.50.21.1:446	134.1.1.99:80
Gateway Received	228.50.21.1:446	134.1.1.99:80
Internal Received	10.18.2.51:446	134.1.1.99:80

- Traditional NAT requires a **one-to-one** relationship between the internal IP address and the external IP address
- Allocating an individual external IP for each internal host's private IP would defeat the purpose of using NAT to preserve IP address space
- One solution was to provide a pool of available external IPs that could be used for NAT connections
- When no more IPs are available, no addition connections to IPs are possible
- With limited backbone IPs available, this would limit the number of IPs available for NAT

- PAT uses the same concept as NAT, except that it rewrites the **port address** as well as the IP address on the external connection
- Port addresses are a 16 bit number, allowing 65534 connections **for a single external IP**
- Domestic ISP Routers utilise PAT to allow multiple concurrent external connections using a single ISP provided external IP

Internal IP	Internal Port	External IP	External Port
10.18.2.20	4673	134.6.101.20	1111
10.18.2.201	11673	134.6.101.20	1112
10.18.2.55	4483	134.6.101.20	1113
10.18.2.222	10987	134.6.101.20	1114

- Extend the life of IPv4 address space
- Provides additional security by hiding internal IP addresses
- Provides independence from ISP provided IPs. Less overhead if change ISP
- Provides flexibility for subnetting both in size and assigned addresses (eg. 10.18.2.51)

- NAT consumes RAM and processor cycles on device performing translation
- Processing time can cause delay in IP communication which may be an issue for time sensitive protocols
- May make it difficult to troubleshoot network issues as TAPs are needed on both sides of the NAT gateway
- Some protocols have issues transitioning across a NAT gateway

- It is possible to have more than one layer of NAT
- This can occur in small-scale ISPs where they have limited external IPs provide customers with an IP from the private IP set. The customer ISP router in turn will provide a NAT gateway for the customer internal network. This can result in significant latency in connections
- The server network may also implement NAT resulting in NAT at both ends of a connection