


ENS1161 Computer Fundamentals

Module 6

Storage

Moving forward..

- ▶ Last module:
 - Programming languages
 - Programming 'tools' that convert programs to binary code
- ▶ Focus of this module: Storage
 - Types of storage
 - How data and instructions are moved from storage to the processor and back



Module Objectives

On completion of this module, students should be able to:

- ▶ Explain the differences between primary and secondary storage.
- ▶ Describe the different types of primary and secondary storage technologies covered and their principles of operation.
- ▶ Explain how the characteristics and specifications of different storage technologies impact their usage in computer systems.
- ▶ Evaluate different types of storage based on their specifications and determine the most suitable for a given application.



Introduction

▶ **Module Scope**

- The role of storage in a computer system
- Types of storage
 - Classification based on usage and technology
- How different storage devices function
- Pros and cons of different storage media


Data & Instructions (recap)

- ▶ Computer systems are **digital** systems
 - All **data** and **instructions** are in the form of **binary signals**
 - **signals** that have only **two** possible values
 - **1** or **0**
 - **HIGH** or **LOW**
 - **ON** or **OFF**
- ▶ These digital signals may be used to represent:
 - one bit of a **binary number**
 - one bit of a **binary code**
 - ASCII, BCD, instruction code, ...
 - a **control signal** state, etc.



Storage (recap)

- ▶ Components used to 'hold' the binary data
 - The **hardware** used to store the **software**
- ▶ **Primary storage / memory**
 - Memory directly accessible by the processor
 - E.g. RAM, ROM, cache memory
 - Fast access, but normally **volatile**
 - Data disappears when power goes off
- ▶ **Secondary storage**
 - Devices that can store data more permanently (even when power off)
 - E.g. hard disk, flash drive, CD/DVD, etc.



Storage

Primary Storage

- Sometimes called **main memory** (or just *memory*)
- Storage that is directly accessible by the CPU
 - Connected directly via **system bus**
- The 'working storage' for the CPU
 - Currently running programs and data stored here
- **Fast access**, but **limited capacity**
- Most of it (RAM) loses data when powered off (**volatile**)

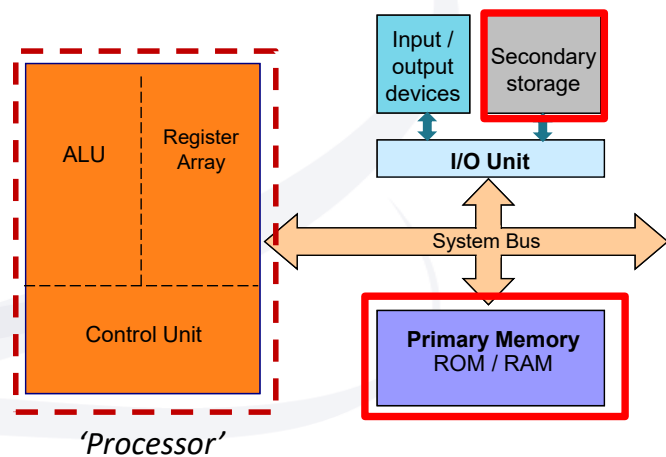
Secondary Storage

- Storage accessed indirectly
 - Connected via **I/O interface**
- Mainly for longer term storage of data
- Can also be used for temporary storage of memory data (*swap space*)
- **Large capacity**, but **slower access**
- Data preserved even when powered off (**non-volatile**)
- Could be physically remote from CPU
- **Cheaper** (lower cost per byte of data)

Basic Components of a Computer (recap - Module 2)

- Every computer contains the same basic components:

- Arithmetic logic unit (ALU)
- Register array
- Control unit
- **Memory**
- **Input/Output** (I/O) unit
- **System Bus**



Storage in a computer





Image: Nick Ares, 2007

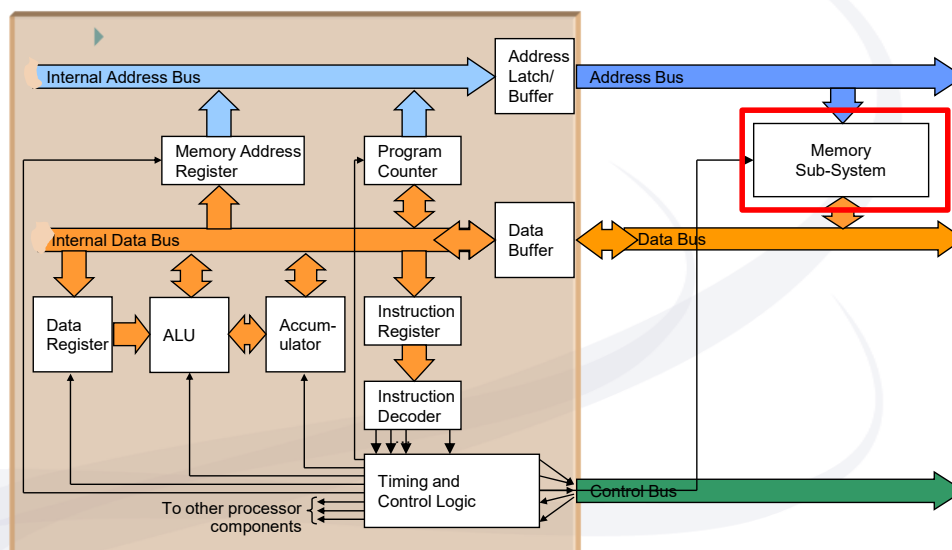
Storage speed vs capacity / frequency of use

- ▶ Trade-off in storage: speed vs capacity
- ▶ Higher speed storage, lower capacity
 - Mainly due to cost per byte
 - Storage with lower access times tend to be more expensive
- ▶ Speed is also normally inversely proportional to distance
 - Further away storage is, the slower it is
 - Due to greater distance, more complex interface
- ▶ Data to be used (more frequently) is transferred to faster ('closer') storage
 - E.g. files stored in secondary storage are transferred to RAM for CPU to access
 - This principle is also used in other mechanisms (*to be covered*)




- Sometimes called **main memory** (or just *memory*)
- Storage that is directly accessible by the CPU
 - Connected directly via [system bus](#)
- The 'working storage' for the CPU
 - Currently running programs and data stored here
- **Fast access**, but **limited capacity**
- Most of it (RAM) loses data when powered off (*volatile*)

Simplified Diagram of a Processor *(recap)*



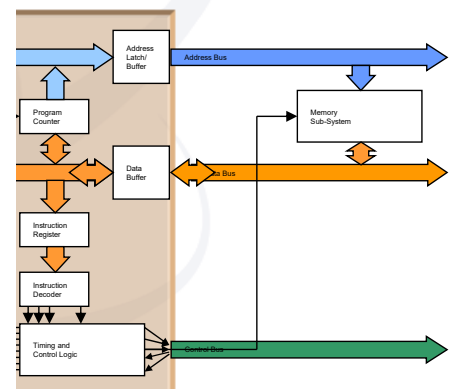
Memory *(recap – Module 2)*

- ▶ Multiple storage locations, each containing binary **data**
 - Generally, each location contains 8 bits (1 byte) of data
 - Could contain multiple bytes (e.g. 16 bits, 32 bits)
- ▶ Each location has an **address**
 - Used to specify location to use
 - Has to be set first so 1 location is selected
- ▶ Data can be read from or written into the selected location
 - Depends on the **control** signal
 - For **read**, data in memory location → data bus
 - For **write**, data on data bus → memory location



Address bus *(recap – Module 2)*

- ▶ Used to select locations within the processors addressable space for reading or writing data
- ▶ The **wider** (more bits) the address bus, the **greater the addressable space**
 - **2^N addresses from an N -bit address bus**
- ▶ The address bus is unidirectional
 - From the processor to memory or device
 - No information is read from it



Data bus (recap - Module 2)

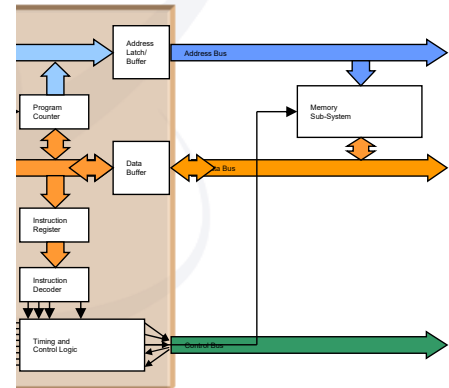
- ▶ Used to transfer data to and from memory (or other device)
- ▶ All information in the system is transferred through this bus
 - data, program instructions, operand addresses, etc.

- ▶ **Width of the data bus** depends on processor

- Normally the number of bits processor can process at one time
- The wider the bus, **more data** can be transferred at one go

- ▶ The data bus is bidirectional

- READ operation: data from *memory location specified on address bus* → processor
- WRITE operation: data from processor → *memory location specified on address bus*




Logic levels

- ▶ **Voltages** are used to represent 0 and 1

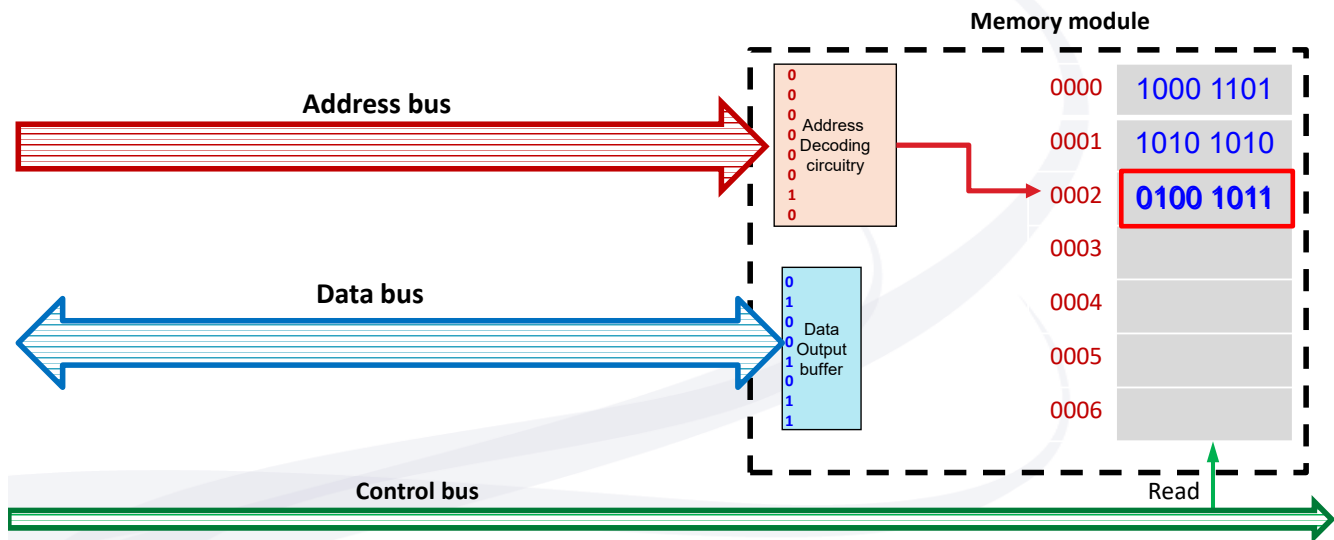
- 0 V nominally used to represent a '0'
- V_{cc} is used to represent '1'
 - V_{cc} = the device power high voltage
 - Typically 5V, but many mobile devices use 3.3V or lower

- ▶ In reality, a **range of voltages** can be accepted as a '1' or '0'

- Range depends on the particular technology used
- $V_{H\ min}$ = lowest voltage that will be read as a '1'
- $V_{L\ max}$ = highest voltage that will be read as a '0'
- In between is an 'indeterminate' range
 - Devices may read them as '1' or '0' (unpredictable)
 - Should not attempt to read when in this range

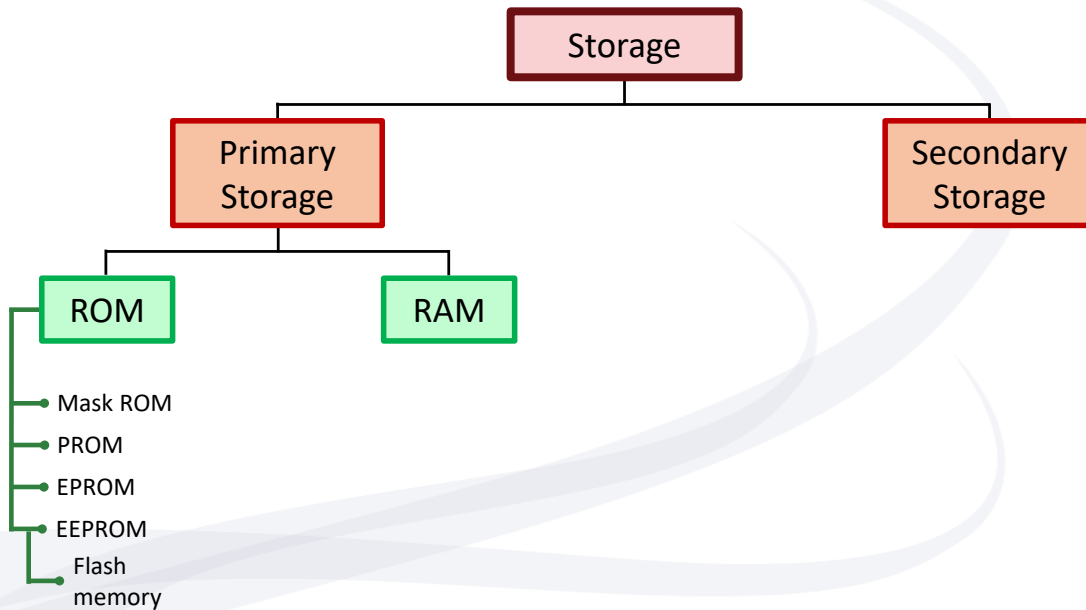
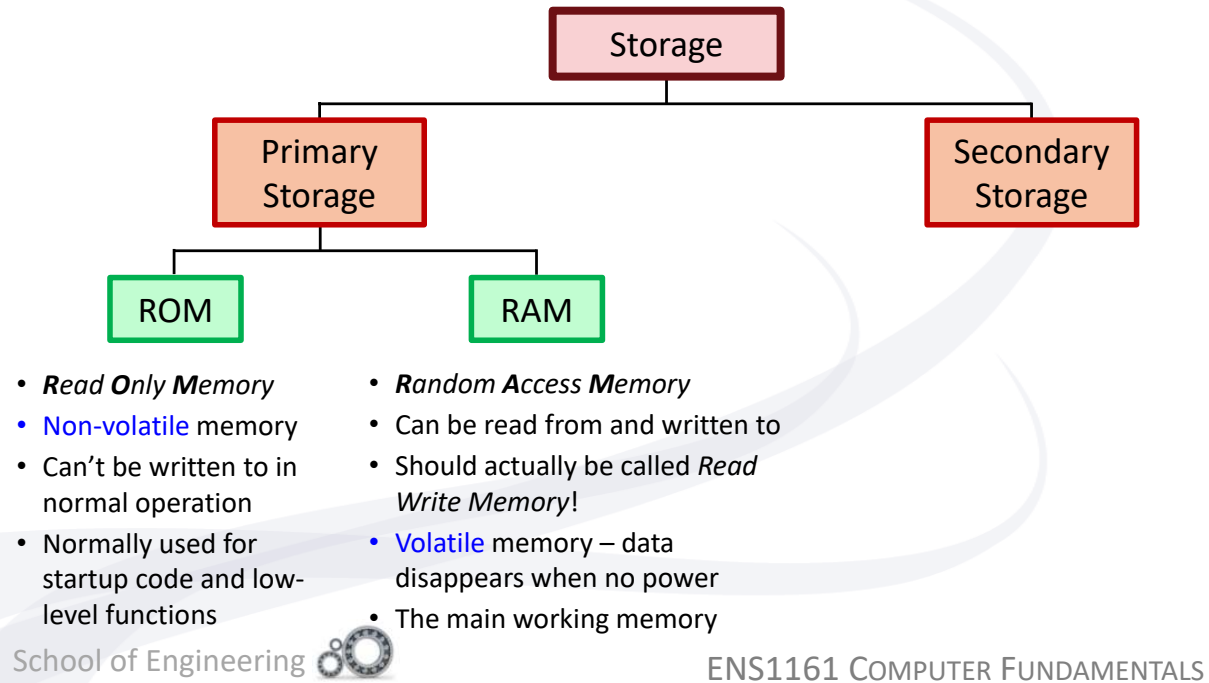


Memory read sequence



Memory access time

- ▶ Memory read takes a **number of steps** from the time the address and read signal are sent till the data is available
 - Allow address signals (voltages) to 'settle' (**stop changing**)
 - **Decode** the address
 - **Select** the correct memory location
 - **Transfer** data to data buffer that is connected to the data bus
 - **Enable** data onto the bus
 - Allow data signals (voltages) to 'settle' (**stop changing**)
- ▶ Each step takes a finite amount of time
- ▶ This is called the **access time** for the memory
 - Normally in the order of 10 – 150 ns (nanoseconds)



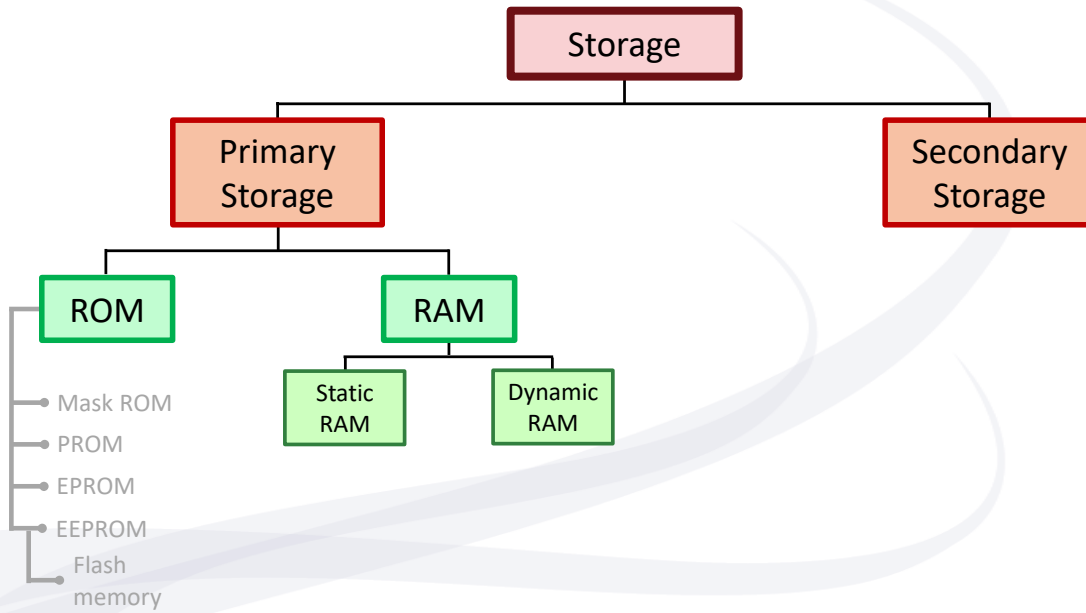
Types of Read Only Memory (ROM)

- **Mask ROMs** – or just ROM
 - Data stored is built into chip at manufacture using a 'mask'
- **PROMs** – (One-time) Programmable ROM
 - Data is written in by 'fusing' (burning) connections
- **EPROM** – Erasable Programmable ROM
 - Data can be cleared by exposing chip to UV light
 - Can then be reprogrammed
- **EEPROM** – Electrical Erasable **PROM**
 - Can be re-programmed *in-circuit*
 - Disadvantage – low density, higher cost
- **Flash Memory** – Higher density **EEPROM**
 - Designed for large block erase and writes



Firmware

- ▶ Software embedded in hardware (ROM)
- ▶ The most common application of ROMs
 - E.g. ROM-**BIOS** (*Basic Input Output System*)
 - Data and program code needed on power-up of computer systems
 - Instructions to initialise the system and invoke an operating system from auxiliary memory
 - *bootstrapping* or '*booting*'
 - Covered in more detail in *Modules 8 and 9*
 - Also low-level functions to handle **I/O** based on hardware in the system
 - **EEPROM** or **Flash Memory** used to allow the **firmware** to be upgraded if necessary




Types of RAM

- ▶ 2 main types of RAM used in computer systems:
 - **Static RAM (SRAM)**
 - Stores data in **logic circuits** (flip-flops)
 - Can stay in a given state indefinitely as long as there is power
 - **Fast** and simple implementation, but **space and power inefficient**
 - **Dynamic RAM (DRAM)**
 - Data stored as charge on small MOS **capacitors**
 - Requires periodic **recharging** - due to **leakage** of capacitors
 - More complex - needs **refresh circuitry**
 - Data read, and rewritten in
 - But can be much **higher density** and **power** requirements are much **lower**

DRAM vs SRAM



DRAM vs. SRAM Memory Cell Complexity



Data stored here

DRAM – principles of operation




- ▶ Essentially just a **capacitor**
 - Acts as a 'bucket' that stores charge
 - Electric current = flow of charge in or out
 - Voltage depends on 'level' of charge in the bucket
 - Whether cell contains '0' or '1' depends on the voltage



DRAM – Limitations

► Time to change data


- 'Bucket' needs time fill up / empty
- Physically impossible to change voltage instantly
- Voltage change follows exponential curve
- This physical limitation applies to all signals
 - Data or addresses
 - Though speed of change depends on circuit
 - 'Size of bucket'



DRAM – Limitations

► Charge leakage

- Capacitors are not perfect
- Charge can leak out
- Voltage will therefore drop
- Data may 'fade away'
- Solution: *Refresh circuitry*
 - Special circuit that periodically checks each cell and 'refreshes' the data
 - Charges it up if a '1' or discharges fully if a '0'
 - 'Refill' or 'empty bucket' accordingly



Processor speed vs RAM speed

- ▶ A single processor instruction may require multiple memory access
 - 1 instruction read (minimum)
 - 1 or more data read or write
 - Depending on type and complexity of instruction
 - *Refer Module 2*
- ▶ Processor clock speeds have increased tremendously
 - E.g. Intel i7 processor standard clock speed = 3.7 GHz
 - i.e. 1 period of clock = 0.27 ns (*nanosecond*)
- ▶ RAM access speeds much slower
 - DRAM typically 60 ns
 - SRAM typically 10 ns
- ▶ While DRAM is cheap and compact, speed can be a bottleneck



Cache Memory

- ▶ Small amount of **faster memory** that sits 'between' CPU and main memory
 - Stores copies of **frequently used** data and instructions
 - To improve performance of primary storage
- ▶ If data required is in cache (**hit**), processor will read data from there
- ▶ If not in cache (**miss**), will then go to main memory to retrieve the data
- ▶ Static RAM used for cache memory
 - Bulkier, more expensive but **faster**



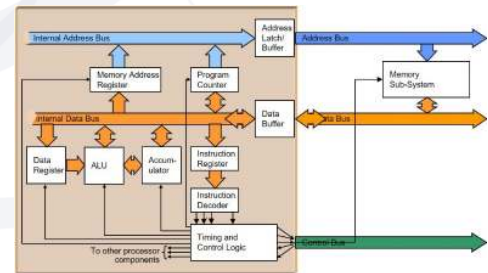


Image: Ferruccio Zulian, Italy




Internal vs External buses (recap Module 2)

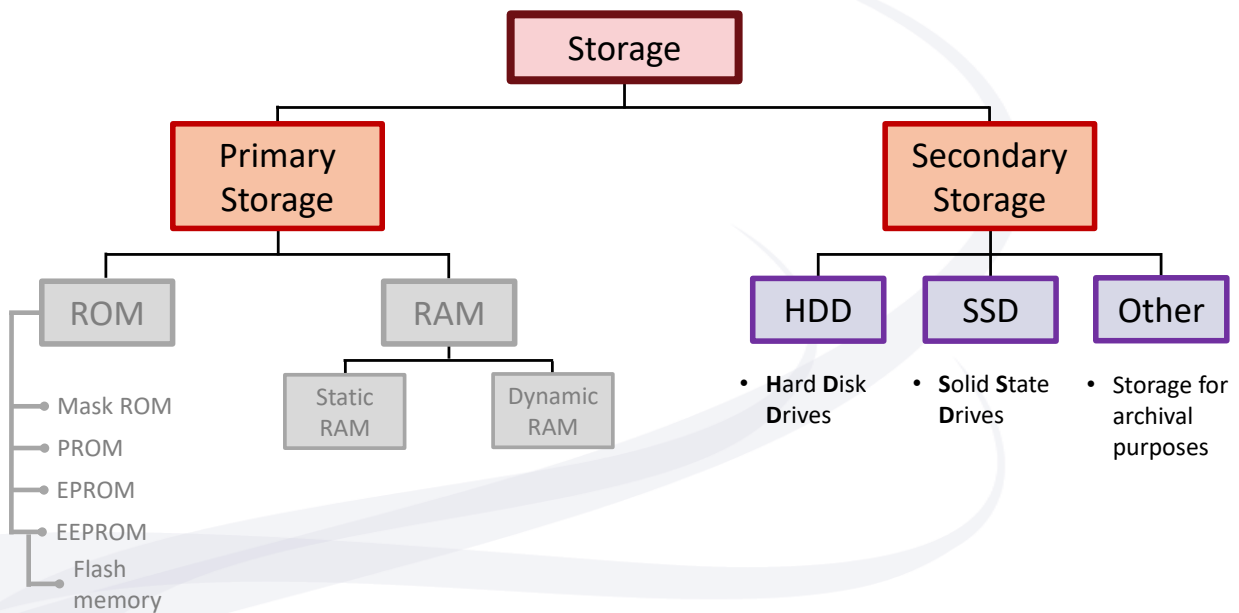
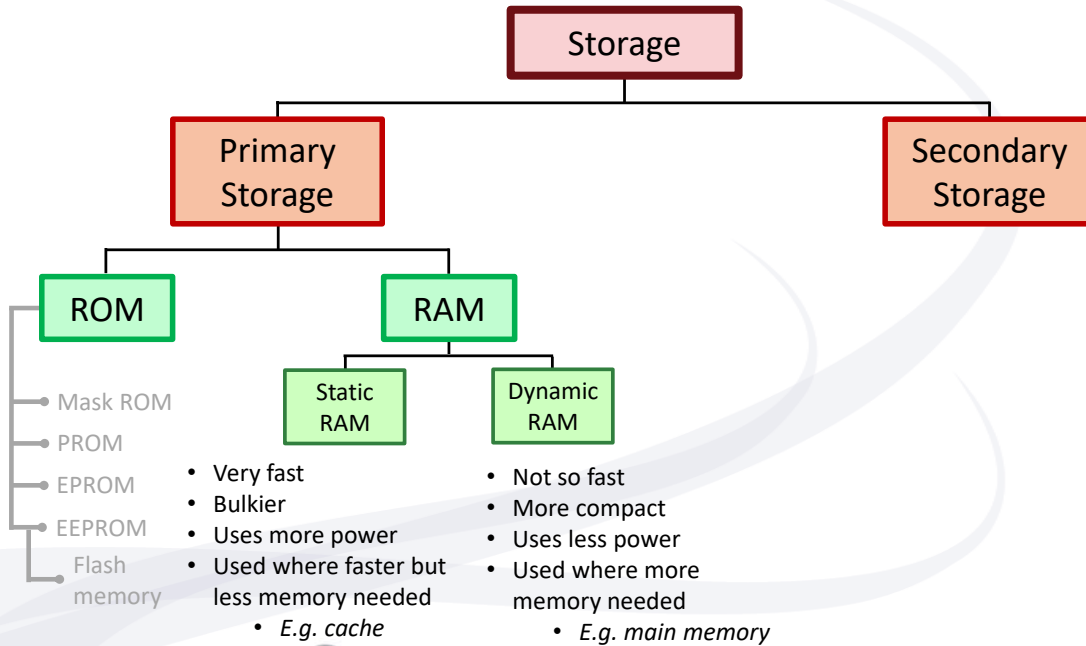
- ▶ The *internal* address and data buses connect the processor components within the processor chip
 - Registers, ALU, etc.
- ▶ The *external* data bus connects the processor chip to external components
 - Memory, I/O devices
- ▶ Internal bus transfers are **much faster** because:
 - **Transfers are within chip**
 - **Shorter distance, less interference**
 - On chip components are faster
 - E.g. registers have a much quicker response time compared to external memory



On-chip Cache


- ▶ **L1 cache** is an *additional level* of cache memory
- ▶ Very fast (but small) cache memory that is *within* the processor chip
 - Takes advantage of the speed of internal transfers within chip
 - Previously described cache memory called *L2 cache*
- ▶ Needed as disparity between processor and memory speeds grew larger
- ▶ Evolving terminology (for multi-core processors):
 - L1 cache: On-chip cache for 1 core
 - L2 cache: On-chip cache shared between cores
 - L3 cache: Off-chip cache






Hard Disk Drives (HDD)

- ▶ Data is stored on metallic **platters** ('hard disks')
 - Platters are coated on both sides with magnetic oxide
 - The magnetic 'orientation' of material changed to represent 1s or 0s
- ▶ **Heads** read the 1's and 0's from the platter
 - Or write them to the platter
 - Heads for top and bottom of each platter
- ▶ Platters spin at high speed
 - Typically 5,400 rpm or higher
- ▶ **Actuator arm** moves the heads together radially
- ▶ Allows any spot to be accessed



HDD data layout

- ▶ Data on hard drives is organised in **tracks**
- ▶ Tracks are subdivided into blocks called **sectors**
 - Normally 512 bytes
- ▶ Operating systems allocate space in **clusters**
 - Groups of sectors
- ▶ Moving the heads from track to track takes time
 - Slows down read / write time
- ▶ So data is normally spread across the same track on the different surfaces of the platters
 - Blocks of data send to different heads
- ▶ This combination of tracks is called a **cylinder**



HDD pros and cons

- ▶ **Access time** is relatively long
 - 4 to 20 ms (*milliseconds*), depending on drive
 - Seek time in addition to normal access delays
 - **Seek time** = time to get head positioned over data location
 - Time to *move head to track*
 - + time for *disk to spin required sector under head*
- ▶ Possibility of **mechanical failure**
 - e.g. *head crash* – head hits platter due to mechanical shock
- ▶ **Power consumption and noise**
 - Because platters always rotating
- ▶ **Cheap**
 - Very low cost per byte
- ▶ **Mature technology**
 - Reliability, speed, noise and power consumption has been improved over time





Image: Wimox, 2019



Disk Fragmentation

- ▶ As files are added and deleted from disk, and files on disk grow and shrink as they are updated, the files become **fragmented**
 - Occupy *non-contiguous* clusters
- ▶ Increases the read time as heads will need to move around from track to track
- ▶ May require **defragmentation** routines to be run to improve performance
 - Move data around on disk so that they lie on contiguous blocks




Image: Enterprise Storage Forum, 2019


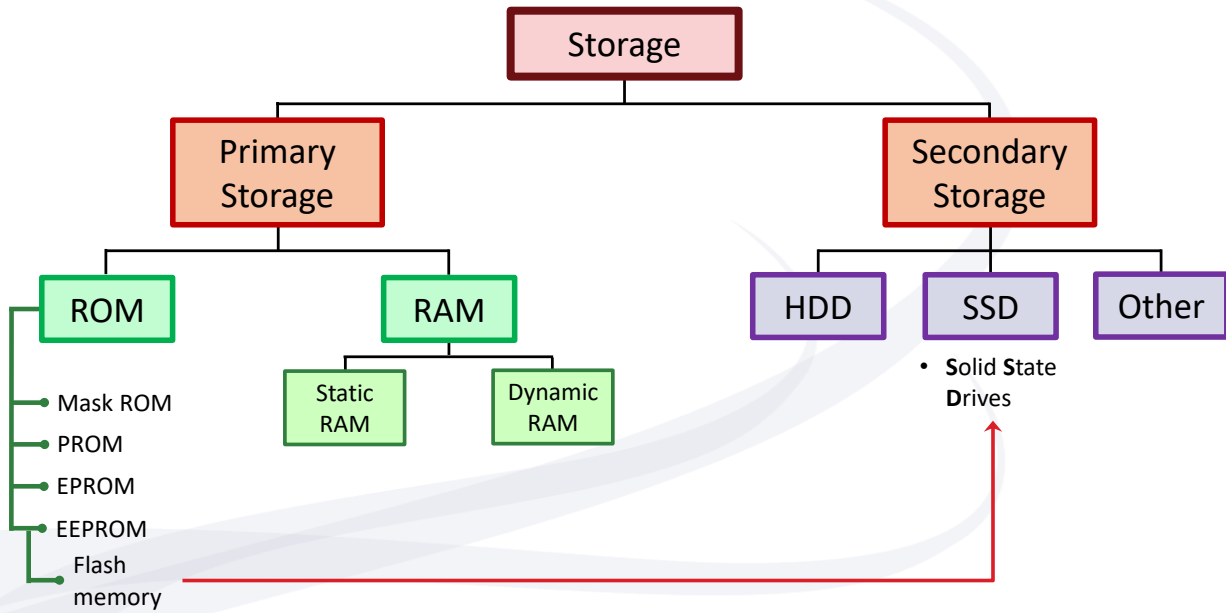


Image: XZise, 2008







Solid State Drives

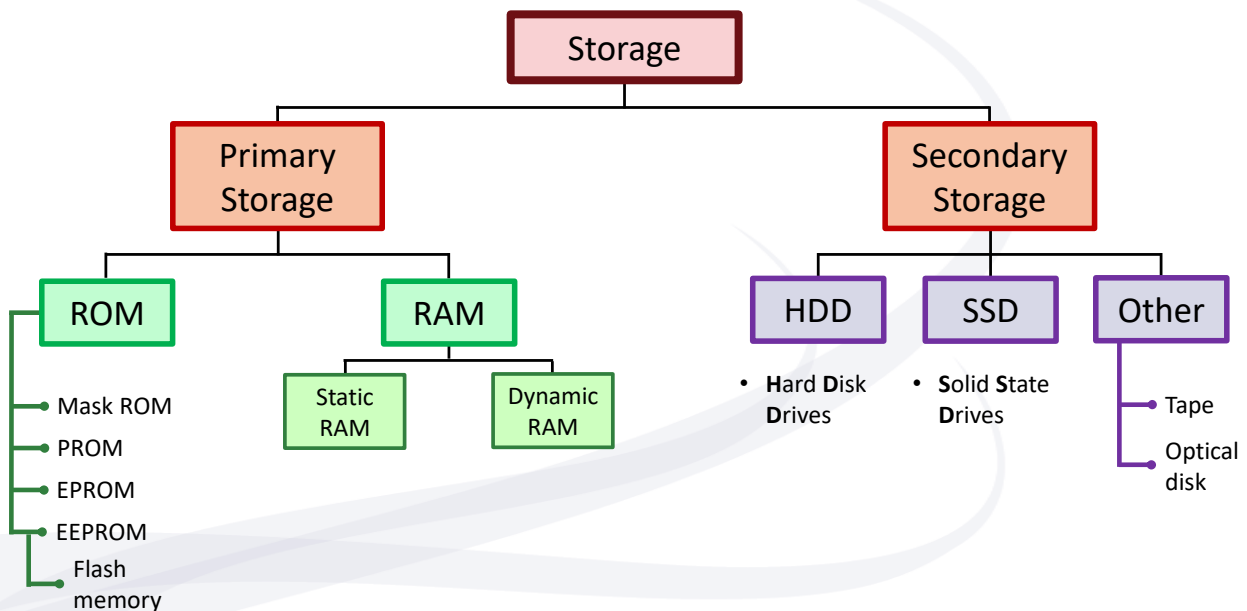
- ▶ Mainly based on Flash ROM technology
 - Works in similar fashion to RAM
 - But transfer blocks of data rather than single bytes
- ▶ Main advantages:
 - **Fast access** (no seek time)
 - Typical access times: 35 – 100 μ s (microseconds)
 - About 100x faster than HDD
 - **Quiet** - no spinning disks
 - **Reliable** - no moving mechanical parts
 - **Fast startup** - no delays waiting for disks to get up to speed
- ▶ Main disadvantage: **High cost** per byte
 - Though price keeps dropping



Image: Ordercrazy, 2014

 SSD vs. HDD Usually 10,000 or 15,000 rpm SAS drives 		
0.1 ms	Access Times SSDs exhibit virtually no access time	5.5-8.0 ms
SSDs deliver at least 6000 io/s	Random I/O Performance SSDs are at least 15 times faster than HDDs	HDDs reach up to 400 io/s
SSDs have a failure rate of less than 0.5%	Reliability This makes SSDs 4-10 times more reliable	HDDs failure rate fluctuates between 2-5%
SSDs consume between 2 and 5 watts	Energy Savings This means that on a large server, approximately 100 watts are saved	HDDs consume between 6 and 15 watts
SSDs have an average I/O wait of 1%	CPU Power You will have an extra 6% of CPU power for other operations	HDDs average I/O wait is about 7%
The average service time for an I/O request while running a backup remain below 20 ms	Input/Output Request Times SSDs allow for much faster data access	The I/O request time with HDDs during backup rises up to 400-500 ms
SSD backups take about 6 hours	Backup Rates SSDs allow for 3-5 times faster backup for your data	HDD backups take up to 20-24 hours

Reproduced with permission.
Copyright 1999-2019 QuinStreet, Inc.
All rights reserved.



Tape

- ▶ Provides sequential access to data
 - Tape needs to be forwarded / rewound to find data
 - Very large access times
- ▶ Earliest computers used tape as secondary storage
- ▶ Now used mainly for **archival** purposes
 - Different formats: DDS-4, Data8, QIC, etc.
- ▶ Cheapest media per byte
- ▶ Tape can safely store data for decades
 - Provided temperature / humidity conditions maintained
 - *c.f. Hard drives*: mechanisms may seize up if not used for years
 - *c.f. Flash*: can lose data if not powered up for 2 – 3 years




Image: Dave Winer, 2008



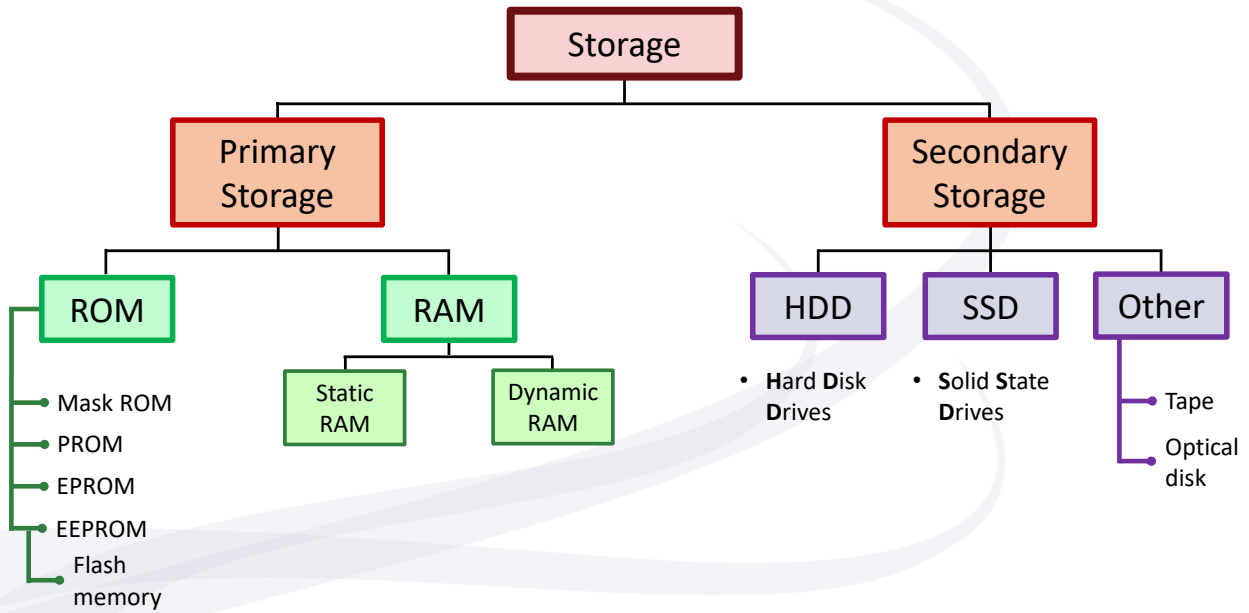
Image: Robert Jacek Tomczak, 2005

Optical discs

- ▶ Data read by laser heads
- ▶ 1s and 0s – ‘lands’ (reflective) and ‘pits’ (non-reflective)
- ▶ Common formats:
 - CD (*Compact Disc*) : 700 Mb
 - DVD (*Digital Versatile Disc*) : 4.7Gb / 9.4 GB
 - Blu-Ray : 25 GB / 50 GB or more
- ▶ Can be read-only, write-once or rewritable
 - Need appropriate disc and drive
- ▶ Physical damage to surface (e.g. scratches) can affect data
- ▶ Usage has declined with the increased popularity of flash drives



Image: By User:Wanted, User:Ochro



RAID

- ▶ Redundant Array of Independent Disks
 - Previously 'Redundant Array of Inexpensive Disks'
- ▶ Have different levels that achieve different objectives
 - Speed
 - To overcome latency of HDD
 - Redundancy
 - Ability to recover data if a drive fails
- ▶ Video: [What is RAID 0, 1, 5, & 10?](#)

Basic RAID Level Summary

- ▶ Raid 0 ('striping')
 - Increases speed by distributing data over different drives
 - No redundancy – worse because if 1 drive fails, all data lost
- ▶ Raid 1 ('mirroring')
 - Data is duplicated over 2 drives
 - Good redundancy – if one drive fails, data still intact
 - Cost – double number of drives
- ▶ Raid 5
 - Has data striped across a number of drives along with *parity* information
 - Data can be rebuilt from parity information if 1 drive fails
- ▶ Raid 10
 - Combines Raid 1 and 0

Module Objectives

On completion of this module, students should be able to:

- ▶ Explain the differences between primary and secondary storage.
- ▶ Describe the different types of primary and secondary storage technologies covered and their principles of operation.
- ▶ Explain how the characteristics and specifications of different storage technologies impacts their usage in computer systems.
- ▶ Evaluate different types of storage based on their specifications and determine the most suitable for a given application.